



NSIGHT COMPUTE

v2022.2.0 | March 2022

Release Notes

TABLE OF CONTENTS

Chapter 1. Release Notes.....	1
1.1. Updates in 2022.2.....	1
1.2. Updates in 2022.1.1.....	2
1.3. Updates in 2022.1.....	2
1.4. Updates in 2021.3.1.....	3
1.5. Updates in 2021.3.....	4
1.6. Updates in 2021.2.4.....	5
1.7. Updates in 2021.2.3.....	5
1.8. Updates in 2021.2.2.....	5
1.9. Updates in 2021.2.1.....	6
1.10. Updates in 2021.2.....	6
1.11. Updates in 2021.1.1.....	7
1.12. Updates in 2021.1.....	8
1.13. Updates in 2020.3.1.....	9
1.14. Updates in 2020.3.....	10
1.15. Updates in 2020.2.1.....	11
1.16. Updates in 2020.2.....	11
1.17. Updates in 2020.1.2.....	12
1.18. Updates in 2020.1.1.....	13
1.19. Updates in 2020.1.....	13
1.20. Updates in 2019.5.3.....	15
1.21. Updates in 2019.5.2.....	15
1.22. Updates in 2019.5.1.....	15
1.23. Updates in 2019.5.....	15
1.24. Updates in 2019.4.....	16
1.25. Updates in 2019.3.1.....	17
1.26. Updates in 2019.3.....	18
1.27. Updates in 2019.2.....	19
1.28. Updates in 2019.1.....	20
Chapter 2. Known Issues.....	22
Chapter 3. Support.....	26
3.1. Platform Support.....	26
3.2. GPU Support.....	27

LIST OF TABLES

Table 1 Platforms supported by NVIDIA Nsight Compute 26

Table 2 GPU architectures supported by NVIDIA Nsight Compute 27

Chapter 1.

RELEASE NOTES

1.1. Updates in 2022.2

General

- ▶ Added support for the CUDA toolkit 11.7.
- ▶ Improved performance for profiling and metric query.
- ▶ Added Linux (aarch64 sbsa) as a supported host platform.
- ▶ The NVIDIA Nsight Compute CLI stores the command line arguments, which can be viewed in the Session report page.
- ▶ Added an API to query the version of the Python Report and NvRules interfaces.
- ▶ Added an API to query the PTX in the Python Report and NvRules interfaces.

NVIDIA Nsight Compute

- ▶ The Acceleration Structure Viewer allows inspection of acceleration structures built using the OptiX API for debugging and performance optimization.
- ▶ The Source page column chooser now supports to enable or disable groups of metrics. Note that not all metrics are enabled anymore by default to make the view easier to use.
- ▶ The Resources tool window now links to the exact target resource instances for CUDA resource types.
- ▶ The Resources tool window now shows the instanced nodes for CUDA graphs.
- ▶ The Resources tool window now shows the loading state and number of loaded functions for CUDA Modules.
- ▶ The Resources tool window now shows the graph node enablement state for applicable instanced graph nodes.
- ▶ The Resources tool window now shows the graph node priorities for instanced kernel graph nodes.
- ▶ Added regex support in the *Next Trigger* filter for NVTX based filtering. The *Next Trigger* filter now considers the NVTX config as a regular expression if the **regex:** prefix is specified.
- ▶ Added regex support in the report's *Filter Results* dialog.

- ▶ Added keyboard shortcuts to navigate between the pages in a report.
- ▶ The behavior for selecting sets and sections is now consistent between the [Sections/Rules Info window](#) and the [non-interactive profile activity](#).

NVIDIA Nsight Compute CLI

- ▶ Added support for collecting [sampling-based](#) warp stalls in range replay mode.
- ▶ Added regex support in [NVTX filtering](#).
- ▶ The metric type is shown when querying metrics.

Resolved Issues

- ▶ Reduced overhead of connecting to the host UI for non-interactive remote profiling sessions.
- ▶ Fixed issues with persisting the Source page state when collapsing or switching between results.
- ▶ Fixed an issue that locked GPU clocks were not reset when terminating the NVIDIA Nsight Compute CLI while profiling a kernel.
- ▶ Fixed issues with selecting and copying text from the Details page tables.
- ▶ Fixed an issue with opening report files in the UI on MacOSX.
- ▶ Fixed an issue with the *Freeze API* option.

1.2. Updates in 2022.1.1

General

- ▶ Filtering kernel launches or profile results based on NVTX domains/ranges now takes registered strings in the payload field into account, if the range name is empty.
- ▶ Added support for the suffix `.max_rate` for ratio metrics.

Resolved Issues

- ▶ Fixed a crash during the disassembly of the kernel's SASS code for the Source page.
- ▶ Fixed a crash on exit of the NVIDIA Nsight Compute UI.
- ▶ Fixed a hang during profiling when CPU call stack collection is enabled.
- ▶ Fixed missing to flush UVM buffers before taking memory checkpoints during [Range Replay](#).
- ▶ Fixed tracking of memory during [Range Replay](#), if the CUDA context has any device mapped memory allocations.
- ▶ Fixed the maximum available shared memory sizes in the [Occupancy Calculator](#) for NVIDIA Ampere GPUs.
- ▶ Fixed that the shared memory usage of the kernel is incorrectly initialized when opening the [Occupancy Calculator](#) from a profile report.

1.3. Updates in 2022.1

General

- ▶ Added support for the CUDA toolkit 11.6.

- ▶ Added support for GA103 chips.
- ▶ Added a new [Range Replay](#) mode to profile ranges of multiple, concurrent kernels. Range replay is available in the NVIDIA Nsight Compute CLI and the non-interactive Profile activity.
- ▶ Added a new rule to detect non-fused floating-point instructions.
- ▶ The Uncoalesced Memory access rules now show results in a dynamic table.
- ▶ Unix Domain Sockets and Windows Named Pipes are used for local connection between the host and target processes on x86_64 Linux and Windows, respectively.
- ▶ The [NvRules API](#) now supports querying action names using different function name bases (e.g. demangled).

NVIDIA Nsight Compute

- ▶ The default [report page](#) is now chosen automatically when opening a report.
- ▶ Added coverage for ECC (Error Correction Code) operations in the L2 Cache table of the Memory Analysis section.
- ▶ Added a new [L2 Evict Policies](#) table to the Memory Analysis section.
- ▶ The [Occupancy Calculator](#) now updates automatically when the input changes.
- ▶ Added new metric *Thread Instructions Executed* to the Source page.
- ▶ Added tooltips to the [Register Dependency](#) columns in the Source page to identify the associated register more conveniently.
- ▶ Improved the selection of Sections and Sets in the Profile activity connection dialog.
- ▶ NVLink utilization is shown in the NVLink Tables section.
- ▶ NVLink links are colored according to the measured throughput.

NVIDIA Nsight Compute CLI

- ▶ **--kernel-regex** and **--kernel-regex-base** options are no longer supported. Alternate options are **--kernel-name** and **--kernel-name-base** respectively, added in 2021.1.0.
- ▶ Added support to resolve CUDA source files in the **--page source** output with the new [--resolve-source-file command line option](#).
- ▶ Added new option **--target-processes-filter** to filter the processes being profiled by name.
- ▶ The CPU Stack Trace is shown in the NVIDIA Nsight Compute CLI output.

Resolved Issues

- ▶ Fixed the calculation of aggregated average instruction execution metrics in non-SASS views on the Source page.
- ▶ Fixed that atomic instructions are counted as both loads and stores in the Memory Analysis tables.

1.4. Updates in 2021.3.1

Resolved Issues

- ▶ Fixed that kernels with the same name and launch configuration were in some scenarios associated with the wrong profiling results during application replay.

- ▶ Fixed an issue with binary forward compatibility of the report format.
- ▶ Fixed an issue with applications calling into the CUDA API during process teardown.
- ▶ Fixed an issue profiling application using pre-CUDA API 3.1 contexts.
- ▶ Fixed a crash when resolving files on the Source page.
- ▶ Fixed that opening reports with large embedded CUBINs would hang the UI.
- ▶ Fixed an issue with remote profiling on a target where the UI is already launched.

1.5. Updates in 2021.3

General

- ▶ Added support for the CUDA toolkit 11.5.
- ▶ Added a new rule for detecting inefficient memory access patterns in the L1TEX cache and L2 cache.
- ▶ Added a new rule for detecting high usage of system or peer memory.
- ▶ Added new `IAction::sass_by_pc` function to the the [NvRules API](#).
- ▶ The [Python-based report interface](#) is now available for Windows and MacOS hosts, too.
- ▶ Added Hierarchical Roofline section files in a new "roofline" section set.
- ▶ Added support for collecting CPU call stack information.

NVIDIA Nsight Compute

- ▶ Added support for new remote profiling [SSH connection and authentication options](#) as well as local SSH configuration files.
- ▶ Added an [Occupancy Calculator](#) which can be opened directly from a profile report or as a new activity. It offers feature parity to the CUDA Occupancy Calculator [spreadsheet](#).
- ▶ Added new [Baselines tool window](#) to manage (hide, update, re-order, save/load) baseline selections.
- ▶ The Source page views now support multi-line/cell selection and copy/paste. Different colors are used for highlighting selections and correlated lines.
- ▶ The search edit on the Source page now supports *Shift+Enter* to search in reverse direction.
- ▶ The [Memory Workload Analysis Chart](#) can be configured to show throughput values instead of transferred bytes.
- ▶ The *Profile* activity now supports the `--devices` option.
- ▶ The *NVLink Topology* diagram displays per NVLink metrics.
- ▶ Added a new tool window showing the CPU call stack at the location where the current thread was suspended during interactive profiling activities.
- ▶ If enabled, the *Call Stack / NVTX* page of the profile report shows the captured CPU call stack for the selected kernel launch.

NVIDIA Nsight Compute CLI

- ▶ Added support for printing source/metric content with the new `--page source` and `--print-source` command line options.

- ▶ Added new option **--call-stack** to enable collecting the CPU call stack for every profiled kernel launch.

Resolved Issues

- ▶ Fixed that **memory_*** metrics could not be collected with the **--metrics** option.
- ▶ Fixed that selection and copy/paste was not supported for section header tables on the Details page.
- ▶ Fixed issues with the Source page when collapsing the content.
- ▶ Fixed that the UI could crash when applying rules to a new profile result.
- ▶ Fixed that PC Sampling metrics were not available for *Profile Series*.
- ▶ Fixed that local profiling did not work if no non-loopback address was configured for the system.
- ▶ Fixed termination of remote-launched applications. On QNX, terminating an application profiled via *Remote Launch* is now supported. Canceling remote-launched *Profile* activities is now supported.

1.6. Updates in 2021.2.4

Resolved Issues

- ▶ Fixed an issue that prevented remote interactive profiling of kernels on NVIDIA GA10b chips.

1.7. Updates in 2021.2.3

General

- ▶ Added support for the NVIDIA GA10b chip.

Resolved Issues

- ▶ Improved error message on QNX for failure to deploy stock section and rules files.

1.8. Updates in 2021.2.2

General

- ▶ Changes for profiling support on NVIDIA virtual GPUs (vGPUs) for an upcoming GRID/vGPU release.

Resolved Issues

- ▶ Fixed hang issue on QNX when using the **--target-processes all** option while profiling shell scripts.

1.9. Updates in 2021.2.1

General

- ▶ Reduced the memory overhead when loading reports in the [Python Report Interface](#).

Resolved Issues

- ▶ Fixed that links in the *Memory Allocations* Resource view were not working correctly.
- ▶ Fixed that NVTX state might not be correctly reset between interactive profiling activities.
- ▶ Fixed that the UI could crash when opening baselines from different GPU architectures.

1.10. Updates in 2021.2

General

- ▶ Added support for the CUDA toolkit 11.4.
- ▶ Added support for OptiX version 7.3.
- ▶ Added support for profiling on [NVIDIA virtual GPUs](#) (vGPUs) on an upcoming GRID/vGPU release.
- ▶ Added a new [Python-based report interface](#) for interacting with report files from Python scripts.
- ▶ Added a new rule to warn users when sampling metrics were selected, but no sampling data was collected.
- ▶ Renamed *SOL* to *Throughput* in the Speed of Light section.
- ▶ Renamed several `memory_*` metrics used on the *Source* page, to better reflect the measured value. See the [Source page](#) documentation for more details.

NVIDIA Nsight Compute

- ▶ Added support for opening [cubin files](#) in a Standalone Source Viewer without profiling the application.
- ▶ Moved the output of all rules so that it is visible even if a section's body is collapsed. Visibility of the rules' output can be toggled by a new button in the report header.
- ▶ The profiler report header now shows the report name for each baseline when ambiguous.
- ▶ Rules can define *Focused Metrics* that were most important for triggering their result output. Metrics are provided per result message which additional information, such as the underlying conditions and thresholds.
- ▶ [Memory tables](#) show tooltips for cells with derived metric calculations.
- ▶ Added a knowledge base service to show more comprehensive background information on metric names and descriptions in their tooltips.
- ▶ Following a link in the Source Counters hot spot tables automatically selects the corresponding metric in the Source page.

- ▶ Added new columns for visualizing register dependencies in the SASS view of the [Source page](#).
- ▶ Functions in the SASS view are now sorted by name.
- ▶ Added support for OptiX 7.x resource tracking in the interactive profile activity. The *Resources* tool window will show information on instantiated `optixDeviceContexts`, `optixModules`, `optixProgramGroups`, `optixPipelines` and `optixDenoiser` objects.
- ▶ Added support for new CUDA graph memory allocation APIs.
- ▶ Improved consistency between command line parameters and the *Next Trigger* filter in the API Stream window for handling of regex inputs. The *Next Trigger* filter now considers kernel/API name as a regular expression only if string has `regex:` as prefix.
- ▶ Added ability to select font settings in the options dialog.
- ▶ Added ability to configure the metrics shown on the summary page via the options dialog.
- ▶ The selected heatmap color scale now also applies to the *Memory chart*.
- ▶ The ncu-ui script now checks for missing library dependencies, such as OpenGL or [Qt](#).

NVIDIA Nsight Compute CLI

- ▶ Added environment variable `NV_COMPUTE_PROFILER_DISABLE_STOCK_FILE_DEPLOYMENT=1` to skip deployment of section and rule files.

Resolved Issues

- ▶ Fixed a performance issue in the NVIDIA Nsight Compute CLI when using `--page raw --csv --units auto`.
- ▶ Fixed that the SSH passphrase key is no longer persisted in the project file.
- ▶ Fixed state of restore button in connection dialog. The button now supports restoring the default settings, if current setting differ from the default.
- ▶ Fixed that the complete GPU name can be shown in the NVLINK topology diagram on MacOS.
- ▶ Fixed that collapsing the Source view reset the selected metrics.
- ▶ Fixed that correlated lines could differ between filtered and unfiltered views of the executed functions.
- ▶ Fixed that two application icons were shown in the MacOS dock.
- ▶ Improved HiDPI awareness.

1.11. Updates in 2021.1.1

General

- ▶ Updated OpenSSL library to version 1.1.1k.

NVIDIA Nsight Compute

- ▶ Remote source resolution can now use the IP address, in addition to the hostname, to find the necessary SSH target.

NVIDIA Nsight Compute CLI

- ▶ Added support for the existing command line options for kernel filtering while importing data from an existing report file using `--import`.
- ▶ Option `-k` is not considered as deprecated option `--kernel-regex` anymore.

Resolved Issues

- ▶ Fixed failure to profile kernels from applications that use the CUDA graphics interop APIs to share semaphores.
- ▶ Fixed wavefront metric in the L1TEX table for writes to shared memory on GA10x chips.
- ▶ Fixed an issue resulting in incomplete data collection for the interactive profile activity after switching from single-pass mode to collecting multiple passes in the same session.
- ▶ Fixed values shown in the mimimap of the Source page when all functions are collapsed.
- ▶ Fixed an issue causing names set by the NVTX naming APIs of one application to be applied to all subsequent sessions of the same instance of NVIDIA Nsight Compute.
- ▶ Fixed behavior of horizontal scroll bars when clicking in the source views on the Source page.
- ▶ Fixed appearance of multi-line entries in column chooser on the Source page.
- ▶ Fixed enablement state of the reset button on the Connection dialog.
- ▶ Fixed potential crash of NVIDIA Nsight Compute when windows size becomes small while being on the Source page.
- ▶ Fixed potential crash of NVIDIA Nsight Compute when relative paths for section/rules files could not be found.
- ▶ Fixed potential crash of NVIDIA Nsight Compute after removing baselines.

1.12. Updates in 2021.1

General

- ▶ Added support for the CUDA toolkit 11.3.
- ▶ Added support for the [OptiX 7 API](#).
- ▶ **GpuArch** enumeration values used for filtering in section files were renamed from architecture names to compute capabilities.
- ▶ NVTX states can now be accessed via the [NvRules API](#).
- ▶ Added a rule for the *Occupancy* section.

NVIDIA Nsight Compute

- ▶ Added support for new CUDA asynchronous allocator attributes in the *Memory Pools* resources view.
- ▶ Added a topology chart and link properties table in the NVLink section.

- ▶ The selected metric column is scrolled into view on the *Source* page when a new metric is selected.
- ▶ Users can choose the *Source* heatmap color scale in the *Options* dialog.

NVIDIA Nsight Compute CLI

- ▶ Added file-based [application replay](#) as the new default application replay mode. File-based replay uses a temporary file for keeping replay data, instead of allocating them in memory. This keeps the required memory footprint close to constant, independent of the number of profiled kernels. Users can switch between buffer modes using the `--app-replay-buffer` option.
- ▶ CLI output now shows NVTX color and message information.
- ▶ `--kernel-regex` and `--kernel-regex-base` options are deprecated and replaced by `--kernel-name` and `--kernel-regex-base`, respectively.
- ▶ All options which support regex need to provide `regex:` as a prefix before an argument to match per the regex, e.g `<option> <regex:expression>`

Resolved Issues

- ▶ Fixed that baselines were not updated properly on the *Comments* page.
- ▶ Fixed that NVTX ranges named using their payloads can be used in [NVTX filtering](#) expressions.
- ▶ Fixed crashes in MacOSX hosts when terminating the target application.
- ▶ The NVLINK(`nv1*`) metrics are now added back.

1.13. Updates in 2020.3.1

General

- ▶ Added support for LDSM instruction-level metrics.

NVIDIA Nsight Compute

- ▶ LDSM instruction-level metrics are shown in the *Source* page and memory tables.
- ▶ Improved reporting and documentation for collecting *Profile Series*.
- ▶ Frozen columns in the *Source* page are automatically scrolled into view.

Resolved Issues

- ▶ Fixed an issue when profiling multi-threaded applications.
- ▶ Fixed an issue that NVIDIA Nsight Compute would not automatically restart when using *Reset Application Data*.
- ▶ Fixed issues with target applications using libstdc++.
- ▶ Fixed an issue when collecting single-pass metrics in multiple Nsight Compute instances.
- ▶ Fixed an issue when using *Kernel ID* and setting *Launch Capture Count* as non-zero in the UI's *Profile* activity.
- ▶ Fixed an issue that prevented different users on the same Linux system to use NVIDIA Nsight Compute in shared instance mode.

- ▶ Fixed an issue that prevented resources from being properly renamed using NVTX information in the UI.

1.14. Updates in 2020.3

General

- ▶ Added support for *derived metrics* in section files. Derived metrics can be used to create new metrics based on existing metrics and constants. See the [Customization Guide](#) for details.
- ▶ Added a new *Import Source* (`--import-source`) option to the UI and command line to permanently import source files into the report, when available.
- ▶ Added a new section that shows selected NVLink metrics on supported systems.
- ▶ Added a new `launch_func_cache_config` metric to the *Launch Statistics* section.
- ▶ Added new branch efficiency metrics to the *Source Counters* section, including `smsp_sass_average_branch_targets_threads_uniform.pct` to replace nvprof's `branch_efficiency`, as well as instruction-level metrics `smsp_branch_targets_threads_divergent`, `smsp_branch_targets_threads_uniform` and `branch_inst_executed`.
- ▶ A warning is shown if kernel replay starts staging GPU memory to CPU memory or the file system.
- ▶ Section and rule files are deployed to a versioned directory in the user's home directory to allow easier editing of those files, and to prevent modifying the base installation.
- ▶ Removed support for NVLINK(`nvl*`) metrics due to a potential application hang during data collection. The metrics will be added back in a future version of the driver/tool.

NVIDIA Nsight Compute

- ▶ Added support for *Profile Series*. Series allow you to profile a kernel with a range of configurable parameters to analyze the performance of each combination.
- ▶ Added a new *Allocations* view to the *Resources* tool window which shows the state of all current memory allocations.
- ▶ Added a new *Memory Pools* view to the *Resources* tool window which shows the state of all current memory pools.
- ▶ Added coverage of peer memory to the *Memory Chart*.
- ▶ The *Source* page now shows the number of excessive sectors requested from L1 or L2, e.g. due to uncoalesced memory accesses.
- ▶ The *Source* column on the *Source* page can now be scrolled horizontally.
- ▶ The kernel duration `gpu_time_duration.sum` was added as column on the *Summary* page.
- ▶ Improved the performance of *application replay* when not all kernels in the application are profiled.

NVIDIA Nsight Compute CLI

- ▶ Added a new `--app-replay-match` option to select the mechanism used for matching kernel instances across application replay passes.
- ▶ An error is shown if `--nvtx-include/exclude` are used without `--nvtx`.

Resolved Issues

- ▶ The *Grid Size* column on the *Raw* page now shows the CUDA grid size like the *Launch Statistics* section, rather than the combined grid and block sizes.
- ▶ The *Branch Resolving* warp stall reason was added to the PC sampling metric groups and the *Warp State Statistics* section.
- ▶ The *API Stream* tool window shows kernel names according to the selected Function Name Mode.
- ▶ Fixed that an incorrect line could be shown after a heatmap selection on the *Source* page.
- ▶ Fixed incorrect metric usage for system memory in the *Memory Chart*. Previously, all requested memory of L2 from system memory was reported instead of only the portion that missed in L2.

1.15. Updates in 2020.2.1

Resolved Issues

- ▶ Fixed several issues related to auto-profiling in the UI.
- ▶ Fixed a metric collection issue when profiling kernels on different GPU architectures with application replay.
- ▶ Fixed a performance problem related to profiling large process trees.
- ▶ Fixed that occupancy charts would not render correctly when comparing against baselines.
- ▶ Fixed that no memory metrics were shown on the *Source* page for `LDGSTS` instructions.
- ▶ Fixed the automatic sorting on the *Summary* and *Raw* pages.
- ▶ Fixed an issue that would cause the NVIDIA Nsight Compute CLI to consume too much memory when importing or printing reports.
- ▶ Long kernel names are now elided in the *Details* page source hot spot tables.
- ▶ Fixed that function names in the *Resources* tool window were demangled differently.

1.16. Updates in 2020.2

General

- ▶ Added support for the NVIDIA Ampere GPUs with compute capability 8.6 and CUDA toolkit 11.1.
- ▶ Added support for application replay to collect metric results across multiple application runs, instead of replaying individual kernels.
- ▶ Added new `launch_device_id` metric.
- ▶ Added support for NVLink (`nvl*`) metrics for GPUs with compute capabilities 7.0, 7.5 and 8.0

- ▶ Added documentation for memory charts and tables in the [Profiling Guide](#).

NVIDIA Nsight Compute

- ▶ Updated menu and toolbar layout.
- ▶ Added support for zoom and pan on roofline charts.
- ▶ The *Resources* tool window shows the current CUDA stream attributes.
- ▶ The memory chart shows a heatmap for link and port utilization.
- ▶ The hot-spot tables in the *Source Counters* section now show values as percentages, too.
- ▶ On-demand resolve of remote CUDA-C source is now available for MacOS hosts.
- ▶ Metric columns in the *Summary* and *Raw* pages are now sortable.
- ▶ Added a new option to set the number of recent API calls shown in the *API Stream* tool window.

NVIDIA Nsight Compute CLI

- ▶ CLI output now shows NVTX payload information.
- ▶ CSV output now shows NVTX states.
- ▶ Added a new **--replay-mode** option to select the mechanism used for replaying a kernel launch multiple times.
- ▶ Added a new **--kill** option to terminate the application once all requested kernels were profiled.
- ▶ Added a new **--log-file** option to decide the output stream for printing tool output.
- ▶ Added a new **--check-exit-code** option to decide if the child application exit code should be checked.

Resolved Issues

- ▶ The profiling progress dialog is not dismissed automatically anymore after an error.
- ▶ The inter-process lock is now automatically given write permissions for all users.
- ▶ All project extensions are enabled in the default dialog filter.
- ▶ Fixed handling of targets using *tcsh* during remote profiling.
- ▶ Fixed handling of quoted application arguments on Windows.

1.17. Updates in 2020.1.2

General

- ▶ The NVIDIA Nsight Compute installer for Mac is now code-signed and notarized.
- ▶ Disabled the creation of the Python cache when executing rules to avoid permission issues and signing conflicts.

Resolved Issues

- ▶ Fixed the launcher script of the NVIDIA Nsight Compute CLI to no longer fail if **uname -p** is not available.
- ▶ Fixed the API parameter capture for function **cuDeviceGetLuid**.

1.18. Updates in 2020.1.1

General

- ▶ Metrics passed to `--metrics` on the NVIDIA Nsight Compute CLI or in the respective *Profile* activity option are automatically expanded to all first-level sub-metrics if required. See the documentation on `--metrics` for more details.
- ▶ Added new rules for detecting inefficiencies of using the sparse data compression on the NVIDIA Ampere architecture.
- ▶ The version of the NVIDIA Nsight Compute target collecting the results is shown in the *Session* page.
- ▶ Added new `launch_grid_dim_[x,y,z]` and `launch_block_dim_[x,y,z]` metrics.

NVIDIA Nsight Compute

- ▶ The *Break on API Error* functionality has been improved when auto profiling.

NVIDIA Nsight Compute CLI

- ▶ The full path to the report output file is printed after profiling.
- ▶ Added and corrected metrics in the nvprof *Metric Comparison* table.

Resolved Issues

- ▶ Documented the *breakdown*: metrics prefix.
- ▶ Fixed handling of escaped domain delimiters in NVTX filter expressions.
- ▶ Fixed issues with the occupancy charts for small block sizes.
- ▶ Fixed an issue when choosing a default report page in the options dialog.
- ▶ Fixed that the scroll bar could overlap the content when exporting the report page as an image.

1.19. Updates in 2020.1

General

- ▶ Added support for the NVIDIA GA100/SM 8.x GPU architecture
- ▶ Removed support for the Pascal SM 6.x GPU architecture
- ▶ Windows 7 is not a supported host or target platform anymore
- ▶ Added a rule for reporting uncoalesced memory accesses as part of the *Source Counters* section
- ▶ Added support for report name placeholders %p, %q, %i and %h
- ▶ The [Kernel Profiling Guide](#) was added to the documentation

NVIDIA Nsight Compute

- ▶ The UI command was renamed from `nv-nsight-cu` to `ncu-ui`. Old names remain for backwards compatibility.
- ▶ Added support for roofline analysis charts

- ▶ Added linked hot spot tables in section bodies to indicate performance problems in the source code
- ▶ Added section navigation links in rule results to quickly jump to the referenced section
- ▶ Added a new option to select how kernel names are shown in the UI
- ▶ Added new memory tables for the L1/TEX cache and the L2 cache. The old tables are still available for backwards compatibility and moved to a new section containing deprecated UI elements.
- ▶ Memory tables now show the metric name as a tooltip
- ▶ Source resolution now takes into account file properties when selecting a file from disk
- ▶ Results in the profile report can now be filtered by NVTX range
- ▶ The Source page now supports collapsing views even for single files
- ▶ The UI shows profiler error messages as dismissible banners for increased visibility
- ▶ Improved the baseline name control in the profiler report header

NVIDIA Nsight Compute CLI

- ▶ The CLI command was renamed from **nv-nsight-cu-cli** to **ncu**. Old names remain for backwards compatibility.
- ▶ Queried metrics on GV100 and newer chips are sorted alphabetically
- ▶ Multiple instances of NVIDIA Nsight Compute CLI can now run concurrently on the same system, e.g. for profiling individual MPI ranks. Profiled kernels are serialized across all processes using a system-wide file lock.

Resolved Issues

- ▶ More C++ kernel names can be properly demangled
- ▶ Fixed a **free() : invalid pointer** error when profiling applications using pytorch > 19.07
- ▶ Fixed profiling IBM Spectrum MPI applications that require PAMI GPU hooks (**--smpiargs="-gpu"**)
- ▶ Fixed that the first kernel instruction was missed when computing **sass_inst_executed_per_opcode**
- ▶ Reduced surplus DRAM write traffic created from flushing caches during kernel replay
- ▶ The *Compute Workload Analysis* section shows the IMMA pipeline on GV11b GPUs
- ▶ Profile reports now scroll properly on MacOS when using a trackpad
- ▶ Relative output filenames for the Profile activity now use the document directory, instead of the current working directory
- ▶ Fixed path expansion of ~ on Windows
- ▶ Memory access information is now shown properly for RED assembly instructions on the Source page
- ▶ Fixed that user **PYTHONHOME** and **PYTHONPATH** environment variables would be picked up by NVIDIA Nsight Compute, resulting in locale encoding issues.

1.20. Updates in 2019.5.3

General

- ▶ More C++ kernel names can be properly demangled

1.21. Updates in 2019.5.2

General

- ▶ Bug fixes

1.22. Updates in 2019.5.1

General

- ▶ Added support for Nsight Compute Visual Studio Integration

1.23. Updates in 2019.5

General

- ▶ Added *section sets* to reduce the default overhead and make it easier to configure metric sets for profiling
- ▶ Reduced the size of the installation
- ▶ Added support for CUDA Graphs Recapture API
- ▶ The NvRules API now supports accessing correlation IDs for instanced metrics
- ▶ Added breakdown tables for *SOL SM* and *SOL Memory* in the Speed Of Light section for Volta+ GPUs

NVIDIA Nsight Compute

- ▶ Added a snap-select feature to the Source page heatmap help navigate large files
- ▶ Added support for loading remote CUDA-C source files via SSH on demand for Linux x86_64 targets
- ▶ Charts on the Details page provide better help in tool tips when hovering metric names
- ▶ Improved the performance of the Source page when scrolling or collapsing
- ▶ The charts for Warp States and Compute pipelines are now sorted by value

NVIDIA Nsight Compute CLI

- ▶ Added support for GPU cache control, see **--cache-control**
- ▶ Added support for setting the kernel name base in command line output, see **--kernel-base**
- ▶ Added support for listing the available names for **--chips**, see **--list-chips**

- ▶ Improved the stability on Windows when using `--target-processes all`
- ▶ Reduced the profiling overhead for small metric sets in applications with many kernels

Resolved Issues

- ▶ Reduced the overhead caused by demangling kernel names multiple times
- ▶ Fixed an issue that kernel names were not demangled in CUDA Graph Nodes resources window
- ▶ The connection dialog better disables unsupported combinations or warns of invalid entries
- ▶ Fixed metric `thread_inst_executed_true` to derive from `smsp_not_predicated_off_thread_inst_executed` on Volta+ GPUs
- ▶ Fixed an issue with computing the theoretical occupancy on GV100
- ▶ Selecting an entry on the Source page heatmap no longer selects the respective source line, to avoid losing the current selection
- ▶ Fixed the current view indicator of the Source page heatmap to be line-accurate
- ▶ Fixed an issue when comparing metrics from Pascal and later architectures on the Summary page
- ▶ Fixed an issue that metrics representing constant values on Volta+ couldn't be collected without non-constant metrics

1.24. Updates in 2019.4

General

- ▶ Added support for the Linux PowerPC target platform
- ▶ Reduced the profiling overhead, especially if no source metrics are collected
- ▶ Reduced the overhead for non-profiled kernels
- ▶ Improved the deployment performance during remote launches
- ▶ Trying to profile on an unsupported GPU now shows an "Unsupported GPU" error message
- ▶ Added support for the `%i` sequential number placeholder to generate unique report file names
- ▶ Added support for `smsp_sass_*` metrics on Volta and newer GPUs
- ▶ The `launch_occupancy_limit_shared_mem` now reports the device block limit if no shared memory is used by the kernel

NVIDIA Nsight Compute

- ▶ The *Profile* activity shows the command line used to launch ncu
- ▶ The heatmap on the Source page now shows the represented metric in its tooltip
- ▶ The *Memory Workload Analysis Chart* on the Details page now supports baselines
- ▶ When applying rules, a message displaying the number of new rule results is shown in the status bar
- ▶ The Visual Profiler Transition Guide was added to the documentation
- ▶ Connection dialog activity options were added to the documentation

- ▶ A warning dialog is shown if the application is resumed without Auto-Profile enabled
- ▶ Pausing the application now has immediate feedback in the toolbar controls
- ▶ Added a *Close All* command to the *File* menu

NVIDIA Nsight Compute CLI

- ▶ The **--query-metrics** option now shows only metric base names for faster metric query. The new option **--query-metrics-mode** can be used to display the valid suffixes for each base metric.
- ▶ Added support for passing response files using the @ operator to specify command line options through a file

Resolved Issues

- ▶ Fixed an issue that reported the wrong executable name in the Session page when attaching
- ▶ Fixed issues that chart labels were shown elided on the Details page
- ▶ Fixed an issue that caused the cache hitrates to be shown incorrectly when baselines were added
- ▶ Fixed an illegal memory access when collecting *sass_*_histogram* metrics for applications using PyTorch on Pascal GPUs
- ▶ Fixed an issue when attempting to collect all *smsp_** metrics on Volta and newer GPUs
- ▶ Fixed an issue when profiling multi-context applications
- ▶ Fixed that profiling start/stop settings from the connection dialog weren't properly passed to the interactive profile activity
- ▶ Fixed that certain *smsp_warp_cycles_per_issue_stall** metrics returned negative values on Pascal GPUs
- ▶ Fixed that metric names were truncated in the **--page details** non-CSV command line output
- ▶ Fixed that the target application could crash if a connection port was used by another application with higher privileges

1.25. Updates in 2019.3.1

NVIDIA Nsight Compute

- ▶ Added ability to send bug reports and suggestions for features using *Send Feedback* in the *Help* menu

Resolved Issues

- ▶ Fixed calculation of theoretical occupancy for grids with blocks that are not a multiple of 32 threads
- ▶ Fixed intercepting child processes launched through Python's subprocess.Popen class
- ▶ Fixed issue of NVTX push/pop ranges not showing up for child threads in NVIDIA Nsight Compute CLI

- ▶ Fixed performance regression for metric lookups on the Source page
- ▶ Fixed description in rule covering the IMC stall reason
- ▶ Fixed cases where baseline values were not correctly calculated in the Memory tables when comparing reports of different architectures
- ▶ Fixed incorrect calculation of baseline values in the Executed Instruction Mix chart
- ▶ Fixed accessing instanced metrics in the NvRules API
- ▶ Fixed a bug that could cause the collection of unnecessary metrics in the Interactive Profile activity
- ▶ Fixed potential crash on exit of the profiled target application
- ▶ Switched underlying metric for **SOL FB** in the GPU Speed Of Light section to be driven by **dram_throughput.avg.pct_of_peak_sustained_elapsed** instead of **fbpa_throughput.avg.pct_of_peak_sustained_elapsed**

1.26. Updates in 2019.3

General

- ▶ Improved performance
- ▶ Bug fixes
- ▶ Kernel launch context and stream are reported as metrics
- ▶ PC sampling configuration options are reported as metrics
- ▶ The default base port for connections to the target changed
- ▶ Section files support multiple, named Body fields
- ▶ NvRules allows users to query metrics using any convertible data type

NVIDIA Nsight Compute

- ▶ Support for filtering kernel launches using their NVTX context
- ▶ Support for new options to select the connection port range
- ▶ The Profile activity supports configuring PC sampling parameters
- ▶ Sections on the Details page support selecting individual bodies

NVIDIA Nsight Compute CLI

- ▶ Support for stepping to kernel launches from specific NVTX contexts
- ▶ Support for new **--port** and **--max-connections** options
- ▶ Support for new **--sampling-*** options to configure PC sampling parameters
- ▶ Section file errors are reported with **--list-sections**
- ▶ A warning is shown if some section files could not be loaded

Resolved Issues

- ▶ Using the **--summary** option works for reports that include invalid metrics
- ▶ The full process executable filename is reported for QNX targets
- ▶ The project system now properly stores the state of opened reports
- ▶ Fixed PTX syntax highlighting
- ▶ Fixed an issue when switching between manual and auto profiling in NVIDIA Nsight Compute

- ▶ The source page in NVIDIA Nsight Compute now works with results from multiple processes
- ▶ Charts on the NVIDIA Nsight Compute details page uses proper localization for numbers
- ▶ NVIDIA Nsight Compute no longer requires the system locale to be set to English

1.27. Updates in 2019.2

General

- ▶ Improved performance
- ▶ Bug fixes
- ▶ Kernel launch context and stream are reported as metrics
- ▶ PC sampling configuration options are reported as metrics
- ▶ The default base port for connections to the target changed
- ▶ Section files support multiple, named Body fields
- ▶ NvRules allows users to query metrics using any convertible data type

NVIDIA Nsight Compute

- ▶ Support for filtering kernel launches using their NVTX context
- ▶ Support for new options to select the connection port range
- ▶ The Profile activity supports configuring PC sampling parameters
- ▶ Sections on the Details page support selecting individual bodies

NVIDIA Nsight Compute CLI

- ▶ Support for stepping to kernel launches from specific NVTX contexts
- ▶ Support for new `--port` and `--max-connections` options
- ▶ Support for new `--sampling-*` options to configure PC sampling parameters
- ▶ Section file errors are reported with `--list-sections`
- ▶ A warning is shown if some section files could not be loaded

Resolved Issues

- ▶ Using the `--summary` option works for reports that include invalid metrics
- ▶ The full process executable filename is reported for QNX targets
- ▶ The project system now properly stores the state of opened reports
- ▶ Fixed PTX syntax highlighting
- ▶ Fixed an issue when switching between manual and auto profiling in NVIDIA Nsight Compute
- ▶ The source page in NVIDIA Nsight Compute now works with results from multiple processes
- ▶ Charts on the NVIDIA Nsight Compute details page uses proper localization for numbers
- ▶ NVIDIA Nsight Compute no longer requires the system locale to be set to English

1.28. Updates in 2019.1

General

- ▶ Support for CUDA 10.1
- ▶ Improved performance
- ▶ Bug fixes
- ▶ Profiling on Volta GPUs now uses the same metric names as on Turing GPUs
- ▶ Section files support descriptions
- ▶ The default sections and rules directory has been renamed to *sections*

NVIDIA Nsight Compute

- ▶ Added new profiling options to the options dialog
- ▶ Details page shows rule result icons in the section headers
- ▶ Section descriptions are shown in the details page and in the sections tool window
- ▶ Source page supports collapsing multiple source files or functions to show aggregated results
- ▶ Source page heatmap color scale has changed
- ▶ Invalid metric results are highlighted in the profiler report
- ▶ Loaded section and rule files can be opened from the sections tool window

NVIDIA Nsight Compute CLI

- ▶ Support for profiling child processes on Linux and Windows x86_64 targets
- ▶ NVIDIA Nsight Compute CLI uses a temporary file if no output file is specified
- ▶ Support for new **--quiet** option
- ▶ Support for setting the GPU clock control mode using new **--clock-control** option
- ▶ Details page output shows the NVTX context when **--nvtx** is enabled
- ▶ Support for filtering kernel launches for profiling based on their NVTX context using new **--nvtx-include** and **--nvtx-exclude** options
- ▶ Added new **--summary** options for aggregating profiling results
- ▶ Added option **--open-in-ui** to open reports collected with NVIDIA Nsight Compute CLI directly in NVIDIA Nsight Compute

Resolved Issues

- ▶ Installation directory scripts use absolute paths
- ▶ OpenACC kernel names are correctly demangled
- ▶ Profile activity report file supports a relative path
- ▶ Source view can resolve all applicable files at once
- ▶ UI font colors are improved
- ▶ Details page layout and label elision issues are resolved
- ▶ Turing metrics are properly reported on the Summary page
- ▶ All byte-based metrics use a factor of 1000 when scaling units to follow SI standards
- ▶ CSV exports properly align columns with empty entries

- ▶ Fixed the metric computation for double_precision_fu_utilization on GV11b
- ▶ Fixed incorrect 'selected' PC sampling counter values
- ▶ The SpeedOfLight section uses 'max' instead of 'avg' cycles metrics for Elapsed Cycles

Chapter 2. KNOWN ISSUES

Installation

- ▶ The Visual Studio 2017 redistributable is not automatically installed by the NVIDIA Nsight Compute installer. The workaround is to install the x64 version of the 'Microsoft Visual C++ Redistributable for Visual Studio 2017' manually. The installer is linked on the main download page for Visual Studio at <https://www.visualstudio.com/downloads/> or download directly from <https://go.microsoft.com/fwlink/?LinkId=746572>.
- ▶ The installer might not show all patch-level version numbers during installation.
- ▶ Some command line options listed in the help of a `.run` installer of NVIDIA Nsight Compute are affecting only the archive extraction, but not the installation stage. To pass command line options to the embedded installer script, specify those options after `--` in the form of `-- <option>`. The available options for the installer script are:

```
-help           : Print help message
-targetpath=<PATH>  : Specify install path
-noprompt       : No prompts. Implies acceptance of the EULA
```

For example, specifying only option `--quiet` extracts the installer archive without any output to the console, but still prompts for user interaction during the installation. To install NVIDIA Nsight Compute without any console output nor any user interaction, please specify `--quiet -- -noprompt`.

Launch and Connection

- ▶ Launching applications on remote targets/platforms is not supported for several combinations. See [Platform Support](#) for details. Manually launch the application using command line `ncu --mode=launch` on the remote system and connect using the UI or CLI afterwards.
- ▶ In the NVIDIA Nsight Compute connection dialog, a remote system can only be specified for one target platform. Remove a connection from its current target platform in order to be able to add it to another.
- ▶ Loading of CUDA sources via SSH requires that the remote connection is configured, and that the hostname/IP address of the connection matches the

- target (as seen in the report session details). For example, prefer my-machine.my-domain.com, instead of my-machine, even though the latter resolves to the same.
- ▶ Other issues concerning remote connections are discussed in the documentation for [remote connections](#).
 - ▶ Local connections between NVIDIA Nsight Compute and the launched target application might not work on some ppc64le or aarch64 (sbsa) systems configured to only support IPv6. On these platforms, the `NV_COMPUTE_PROFILER_LOCAL_CONNECTION_OVERRIDE=uds` environment variable can be set to use *Unix Domain Sockets* instead of TCP for local connections to workaround the problem. On x86_64 Linux, Unix Domain Sockets are used by default but local TCP connections can be forced using `NV_COMPUTE_PROFILER_LOCAL_CONNECTION_OVERRIDE=tcp`.

Profiling and Metrics

- ▶ Profiling kernels executed on a device that is part of an SLI group is not supported. An "Unsupported GPU" error is shown in this case.
- ▶ Profiling a kernel while other contexts are active on the same device (e.g. X server, or secondary CUDA or graphics application) can result in varying metric values for L2/FB (Device Memory) related metrics. Specifically, L2/FB traffic from non-profiled contexts cannot be excluded from the metric results. To completely avoid this issue, profile the application on a GPU without secondary contexts accessing the same device (e.g. no X server on Linux).
- ▶ In the current release, profiling a kernel while any other GPU work is executing on the same MIG compute instance can result in varying metric values for all units. NVIDIA Nsight Compute enforces serialization of the CUDA launches within the target application to ensure those kernels do not influence each other. See [Serialization](#) for more details. However, GPU work issued through other APIs in the target process or workloads created by non-target processes running simultaneously in the same MIG compute instance will influence the collected metrics. Note that it is acceptable to run CUDA processes in other MIG compute instances as they will not influence the profiled MIG compute instance.
- ▶ On Linux kernels settings `fs.protected_regular=1` (e.g. some Ubuntu 20.04 cloud service provider instances), root users may not be able to access the [inter-process lock file](#). See the [FAQ](#) for workarounds.
- ▶ Profiling only supports up to 32 device instances, including instances of MIG partitions. Profiling the 33rd or higher device instance will result in indeterminate data.
- ▶ Enabling certain metrics can cause GPU kernels to run longer than the driver's watchdog time-out limit. In these cases the driver will terminate the GPU kernel resulting in an application error and profiling data will not be available. Please disable the driver watchdog time out before profiling such long running CUDA kernels.
 - ▶ On Linux, setting the X Config option Interactive to false is recommended.
 - ▶ For Windows, detailed information on disabling the Windows TDR is available at <https://docs.microsoft.com/en-us/windows-hardware/drivers/display/timeout-detection-and-recovery>

- ▶ Collecting device-level metrics, such as the NVLINK metrics (`nv1*`), is not supported on [NVIDIA virtual GPUs](#) (vGPUs).
- ▶ As of CUDA 11.4 and R470 TRD1 driver release, NVIDIA Nsight Compute is supported in a vGPU environment which requires a vGPU license. If the license is not obtained after 20 minutes, the reported performance metrics data from the GPU will be inaccurate. This is because of a feature in vGPU environment which reduces performance but retains functionality as specified [here](#).
- ▶ Profiling on [NVIDIA live-migrated virtual machines](#) is not supported and can result in undefined behavior.
- ▶ Profiling with enabled multi-process service (MPS) can result in undefined behavior.
- ▶ The NVLink Topology section is not supported for a configuration using NVSwitch.
- ▶ NVIDIA Nsight Compute does not support per-NVLink metrics.
- ▶ NVIDIA Nsight Compute does not support the *Logical NVLink Throughput* table.

Compatibility

- ▶ Reports collected on Windows might show invalid characters for file and process names when opened in NVIDIA Nsight Compute on Linux.
- ▶ Applications calling blocking functions on std input/output streams can result in the profiler to stop, until the blocking function call is resolved.
- ▶ NVIDIA Nsight Compute can hang on applications using RAPIDS in versions 0.6 and 0.7, due to an issue in cuDF.
- ▶ Profiling child processes launched via `clone()` is not supported.
- ▶ Profiling child processes launched from Python using `os.system()` is not supported.
- ▶ Profiling of Cooperative Groups kernels launched with `cuLaunchCooperativeKernelMultiDevice` is not yet supported.
- ▶ On Linux systems, when profiling *bsd-csh* scripts, the original application output will not be printed. As a workaround, use a different C-shell, e.g. *tcs*.
- ▶ Attempting to use the `--clock-control` option to set the GPU clocks will fail when profiling on a GPU partition. Please use `nvidia-smi` (installed with nvidia display driver) to control the clocks for the entire GPU. This will require administrative privileges when the GPU is partitioned.
- ▶ On Linux aarch64, NVIDIA Nsight Compute does not work if the *HOME* environment variable is not set.
- ▶ NVIDIA Nsight Compute versions 2020.1.0 to 2020.2.1 are not compatible with CUDA driver version 460+ if the application launches Cooperative Groups kernels. Profiling will fail with error "UnknownError".
- ▶ Collecting CPU call stack information on Windows Server 2016 can hang NVIDIA Nsight Compute in some cases. Currently, the only workaround is to skip CPU call stack collection on such systems by not specifying the option `--call-stack`.

User Interface

- ▶ The API Statistics filter in NVIDIA Nsight Compute does not support units.
- ▶ File size is the only property considered when resolving source files. Timestamps are currently ignored.

- ▶ Terminating or disconnecting an application in the *Interactive Profiling* activity while the API Stream View is updated can lead to a crash.
- ▶ See the [OptiX library support section](#) for limitations concerning the [Acceleration Structure Viewer](#).

Chapter 3. SUPPORT

Information on supported platforms and GPUs.

3.1. Platform Support

Host denotes the UI can run on that platform. Target means that we can instrument applications on that platform for data collection. Applications launched with instrumentation on a target system can be connected to from most host platforms. The reports collected on one system can be opened on any other system.

Table 1 Platforms supported by NVIDIA Nsight Compute

	Host	Targets
Windows	Yes	Windows*, Linux (x86_64)
Windows Subsystem for Linux	No	No
Linux (x86_64)	Yes	Windows*, Linux (x86_64), Linux (ppc64le), Linux (aarch64 sbsa)
Linux (ppc64le)	No	Linux (ppc64le)
Linux (aarch64 sbsa)	Yes	Linux (aarch64 sbsa)
Linux (x86_64) (Drive SDK)	Yes	Windows*, Linux (x86_64), Linux (aarch64), QNX
MacOSX 10.13+	Yes	Windows*, Linux (x86_64), Linux (ppc64le)
Linux (aarch64)	No	Linux (aarch64)
QNX	No	QNX

Target platforms marked with * do not support remote launch from the respective host. Remote launch means that the application can be launched on the target system from the host UI. Instead, the application must be launched from the target system.

On all Linux platforms, NVIDIA Nsight Compute CLI requires GLIBC version 2.15 or higher. On x86_64, the NVIDIA Nsight Compute UI requires GLIBC version 2.17

or higher. On Linux (aarch64 sbsa), the NVIDIA Nsight Compute UI requires GLIBC version 2.26 or higher.

Only Windows 10 and 11 are supported as host and target.

Profiling of 32-bit processes is not supported.

3.2. GPU Support

Table 2 GPU architectures supported by NVIDIA Nsight Compute

Architecture	Support
Kepler	No
Maxwell	No
Pascal	No
Volta GV100	Yes
Volta GV11b	Yes
Turing TU10x	Yes
NVIDIA GA100	Yes
NVIDIA GA10x	Yes
NVIDIA GA10b	Yes

Most metrics used in NVIDIA Nsight Compute are identical to those of the PerfWorks Metrics API and follow the documented [Metrics Structure](#). A comparison between the metrics used in nvprof and their equivalent in NVIDIA Nsight Compute can be found in the [NVIDIA Nsight Compute CLI User Manual](#).

Notice

ALL NVIDIA DESIGN SPECIFICATIONS, REFERENCE BOARDS, FILES, DRAWINGS, DIAGNOSTICS, LISTS, AND OTHER DOCUMENTS (TOGETHER AND SEPARATELY, "MATERIALS") ARE BEING PROVIDED "AS IS." NVIDIA MAKES NO WARRANTIES, EXPRESSED, IMPLIED, STATUTORY, OR OTHERWISE WITH RESPECT TO THE MATERIALS, AND EXPRESSLY DISCLAIMS ALL IMPLIED WARRANTIES OF NONINFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE.

Information furnished is believed to be accurate and reliable. However, NVIDIA Corporation assumes no responsibility for the consequences of use of such information or for any infringement of patents or other rights of third parties that may result from its use. No license is granted by implication of otherwise under any patent rights of NVIDIA Corporation. Specifications mentioned in this publication are subject to change without notice. This publication supersedes and replaces all other information previously supplied. NVIDIA Corporation products are not authorized as critical components in life support devices or systems without express written approval of NVIDIA Corporation.

Trademarks

NVIDIA and the NVIDIA logo are trademarks or registered trademarks of NVIDIA Corporation in the U.S. and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

Copyright

© 2018-2022 NVIDIA Corporation and affiliates. All rights reserved.

This product includes software developed by the Syncro Soft SRL (<http://www.sync.ro/>).